

Acceptor and Connector

Design Patterns for Initializing Communication Services

Douglas C. Schmidt
schmidt@cs.wustl.edu

Department of Computer Science
Washington University
St. Louis, MO 63130, USA
(314) 935-7538

This paper will appear at the EuroPLoP '96 Conference in Kloster Irsee, Germany.

1 Introduction

This paper describes the Connector and Acceptor patterns. The intent of these patterns is to decouple the active and passive connection roles, respectively, from the tasks a communication service performs once connections are established. Common examples of communication services that utilize these patterns include WWW browsers, WWW servers, object request brokers, and “superservers” that provide services like remote login and file transfer to client applications.

This paper illustrates how the Connector and Acceptor patterns can help decouple the connection-related processing from the service processing, thereby yielding more reusable, extensible, and efficient communication software. When used in conjunction with related patterns like the Reactor [1], Active Object [2], and Service Configurator [3], the Acceptor and Connector patterns enable the creation of highly extensible and efficient communication software frameworks [4] and applications [5].

This paper is organized as follows: Section 2 outlines background information on networking and communication protocols necessary to appreciate the patterns in this paper; Section 3 motivates the need for the Acceptor and Connector patterns and illustrates how they have been applied to a production application-level Gateway; Section 4 describes the Acceptor and Connector patterns in detail; and Section 5 presents concluding remarks.

2 Background

Connection-oriented protocols (such as TCP [6]) reliably deliver data between two or more endpoints of communication. Initializing these endpoints involves the following two roles:

- *The passive role* – which initializes an endpoint of communication at a particular address and waits passively for the other endpoint(s) to connect with it;

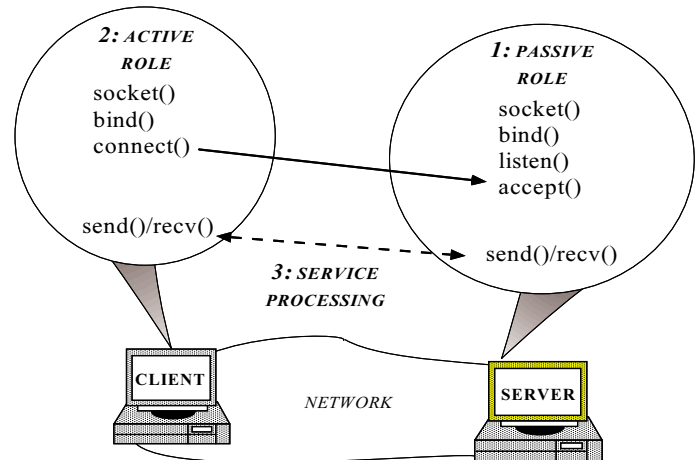


Figure 1: Active and Passive Initialization Roles

- *The active role* – which actively initiates a connection to one or more endpoints that are playing the passive role.

Figure 1 illustrates how these initialization roles behave and interact when a client actively connects to a passive server using the socket network programming interface [7] and the TCP transport protocol [8]. In this figure the server plays the passive initialization role and the client plays the active initialization role.¹

The primary goal of the Acceptor and Connector patterns is to decouple the passive and active initialization roles, respectively, from the tasks performed once the endpoints of a service are initialized. These patterns are motivated by the observation that the tasks performed on messages exchanged between endpoints of a distributed service are largely independent of the following:

- **Which endpoint initiated the connection:** Connection establishment is inherently asymmetrical since the passive endpoint *waits* and the active endpoint *initiates* the connection. Once the connection is established, however, data may be transferred between endpoints in any manner that

¹The distinction between “client” and “server” refer to *communication* roles, not necessarily to *connection* roles. Although clients often play the active role when initiating connections with a passive server these initialization roles can be reversed, as shown in Section 3.

obeys the service's communication protocol (e.g., peer-to-peer, request-response, oneway streaming, etc.). Figure 1 illustrates two endpoints exchanging messages once a connection is established and the client-side and server-side of the service is initialized.

- **The network programming interfaces and underlying protocols used to establish the connection:** Different network programming interfaces (such as sockets [9] or TLI [10]) provide different library calls to establish connections using various underlying communication protocols (such as the Internet TCP/IP protocol or Novell's IPX/SPX). Regardless of the mechanism used to establish a connection, however, data can be transferred between endpoints using uniform message passing operations (e.g., UNIX read/write calls or Win32 ReadFile/WriteFile).

- **The creation, connection, and concurrency strategies used to initialize and execute the service:** The processing tasks performed by a service are often independent of the strategies used (1) to create an instance of a service, (2) connect the service instance to one or more peers, and (3) execute this service instance in one or more threads or processes. By explicitly decoupling these initialization strategies from the behavior of the service, the Connector and Acceptor patterns increase the potential for reusing and extending the service in different environments.

3 Motivation

3.1 Context

To illustrate the Acceptor and Connector patterns, consider the multi-service, application-level Gateway shown in Figure 2. The Gateway routes several types of data (such as status information, bulk data, and commands) that are exchanged between services running on the Peers. The Peers are used to monitor and control a satellite constellation. They can be located throughout local area networks (LANs) and wide-area networks (WANs).

The Gateway is a Mediator [11] that coordinates interactions between its connected Peers. From the Gateway's perspective, these Peer services differ solely by their message framing formats and payload types. The Gateway uses a connection-oriented interprocess communication (IPC) mechanism (such as TCP) to transmit data between its connected Peers. Using a connection-oriented protocol simplifies application error handling and enhances performance over long-latency WANs.

Each communication service in the Peers sends and receives status information, bulk data, and commands to and from the Gateway using separate TCP connections. Each connection is bound to a unique address (e.g., an IP address and port number). For example, bulk data sent from a ground station Peer through the Gateway is connected to a different port than status information sent by a tracking station peer through the Gateway to a ground station Peer. Separating connections in this manner allows more flexible routing

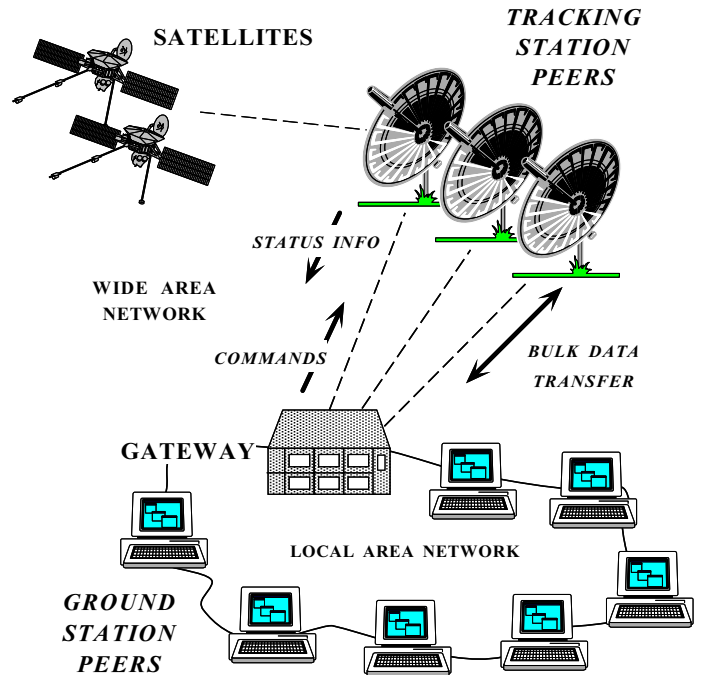


Figure 2: The Physical Architecture of a Connection-oriented Application-level Gateway

strategies and more robust error handling when connections fail.

3.2 Common Traps and Pitfalls

One way to design the Peers and Gateway is to designate the connection roles *a priori*. For instance, the Gateway could be hard-coded to actively initiate the connections for all its services. To accomplish this, it could iterate through a list of Peers and synchronously connect with each of them. Likewise, Peers could be hard-coded to passively accept connections and initialize the associated services. Moreover, the active and passive connection code for the Gateway and Peers, respectively, could be implemented with conventional network programming interfaces (such as sockets or TLI). In this case, a Peer could call `socket`, `bind`, `listen`, and `accept` to initialize a passive-mode listener socket and the Gateway could call `socket` and `connect` to actively initiate a data-mode connection socket. Once the connections were established, the Gateway could route data for each type of service it provided.

However, this approach has several drawbacks:

- **Limited extensibility and reuse of the Gateway and Peer software:** The type of routing service (e.g., status information, bulk data, or commands) performed by the Gateway is independent of the mechanisms used to establish connections and initialize services. However, the approach described above tightly couples the service initialization and service behavior. This makes it hard to reuse existing services or

to extend the Gateway by adding new routing services and enhancing existing services.

- **Error-prone network programming interfaces:** conventional network programming interfaces (such as sockets or TLI) do not provide adequate type-checking since they utilize low-level I/O handles [12]. The tight coupling of the approach describe above makes it easy to accidentally misuse these interfaces and I/O handles in ways that cannot be detected until run-time.
- **Lack of scalability:** If there are a large number of Peers, the synchronous connection establishment strategy of the Gateway will not take advantage of the parallelism inherent in the network and Peers.

3.3 Solution

A more flexible and efficient way to design the Peers and Gateway is to use the *Acceptor* and *Connector* patterns. These two patterns decouple the *active* and *passive* initialization roles, respectively, from the communication services performed once services are initialized. These patterns resolve the following forces for communication services that use connection-oriented transport protocols:

1. **The need to avoid rewriting initialization code for each new service:** The Connector and Acceptor patterns permit key characteristics of services (such as application-level communication protocols and data formats) to evolve independently of the strategies used to initialize the services. These service characteristics often change more frequently than initialization strategies. Therefore, this separation of concerns helps reduce software coupling and increases code reuse.
2. **The need to make connection establishment software portable across platforms:** Many operating systems provide network programming interfaces (such as sockets and TLI) and communication protocols (such as TCP/IP and IPX/SPX) whose semantics are only superficially different. Therefore, the syntactic incompatibilities of these interfaces make it hard to write portable programs, even though the core initialization strategies transcend these differences. It is particularly hard to write portable asynchronous connection establishment software since asynchrony is not supported uniformly by many network programming interfaces.
3. **The need to enable flexible strategies for executing communication services concurrently:** Once a connection is established, peer applications use the connection to exchange data to perform some type of service, such as a network time service, a WWW HTML document transfer service, or a remote login service. However, regardless of how the connection was established, these services may be executed in a single-thread, in multiple threads, or multiple processes, respectively.

4. **The need to actively establish connections with large number of peers efficiently:** The Connector pattern can employ asynchrony to initiate and complete multiple connections without blocking the caller. By using asynchrony, the Connector pattern enables applications to actively establish connections with a large number of peers efficiently over long-latency WANs.

5. **The need to ensure that passive-mode I/O handles are not accidentally used to read or write data:** Strongly decoupling the Acceptor initialization role from the role of the initialized service ensures that passive-mode listener endpoints are not accidentally used incorrectly. Without this strong decoupling, services may mistakenly read or write data on passive-mode listener endpoints, which should only be used to accept connections.

As outlined above, the Connector and Acceptor patterns address very similar forces. For instance, forces 1, 2, and 3 above are resolved by both the Acceptor and Connector pattern – only the passive and active roles are reversed. Forces 4 and 5 are not resolved by both patterns, however, due to the asymmetrical connection roles played by each pattern. For example, the Connector pattern addresses an additional force (connection scalability) by using asynchrony to actively establish connections with a large number of peers efficiently. This force need not be addressed by the Acceptor since it is always the passive target of active connection requests. Conversely, a Connector does not wait passively for services to connect with it, so this pattern need not decouple the listener endpoint from the data endpoint, as the Acceptor does.

The following section describes the Acceptor and Connector patterns using pattern form.

4 The Acceptor and Connector Patterns

4.1 Intent

The intent of these patterns is to decouple service initialization from the tasks performed once a service is initialized. The Connector pattern is responsible for *active* initialization, whereas the Acceptor pattern is responsible for *passive* initialization.

4.2 Also Known As

The Acceptor pattern is also known as the Listener [10].

4.3 Applicability

- Use the Acceptor and Connector patterns when tasks performed by a service can be decoupled from the steps required to initialize a service;
- Use the Connector pattern when an application must establish a large number of connections with peers con-

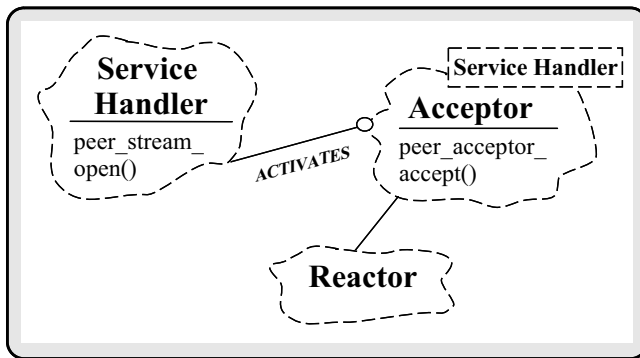


Figure 3: Structure of Participants in the Acceptor Pattern

ected over long-latency networks (such as satellite WANs);

- Use the Acceptor pattern when connections may arrive concurrently from different peers, but blocking or continuous polling for incoming connections on any individual peer is inefficient.

4.4 Structure and Participants

The structure of the participants in the Acceptor and Connector patterns are illustrated by the Booch class diagram [13] in Figure 3 and Figure 4, respectively.² The participants that are common to each pattern are described first, followed by participants that are unique to the Acceptor and Connector patterns:

- **Reactor**

- For the Acceptor pattern, the Reactor demultiplexes connection requests received on one or more communication endpoints to the appropriate Acceptor (described below). The Reactor allows multiple Acceptors to listen for connections from peers within a single thread of control. For the Connector pattern, the Reactor handles the completion of connections that were initialized asynchronously. The Reactor allows multiple Service Handlers to have their connections initiated and completed asynchronously by a Connector configured within a single thread of control.

- **Service Handler**

- Defines a generic interface for a service. The Service Handler contains a communication endpoint (`peer_stream`) that encapsulates an I/O handle (also known as an “I/O descriptor”). This endpoint is initialized by the

²In these diagrams dashed clouds indicate classes; dashed boxes in the clouds indicate template parameters; a solid undirected edge with a hollow circle at one end indicates a uses relation between two classes.

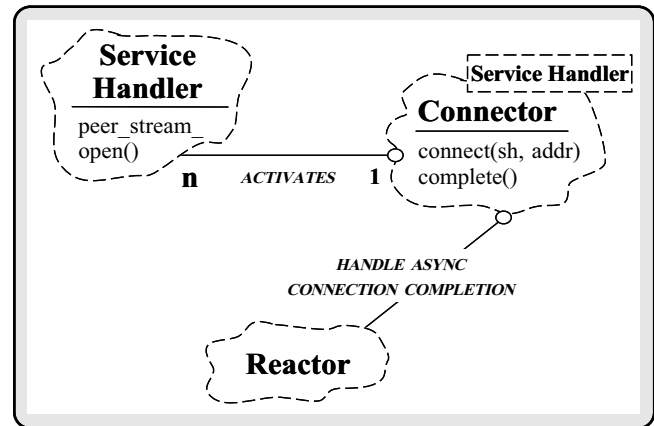


Figure 4: Structure of Participants in the Connector Pattern

Acceptor and Connector and is subsequently used by the Service Handler to exchange data with its connected peer. The Acceptor and Connector activate a Service Handler by calling its open method when a connection is established. Once a Service Handler is completely initialized (by either an Acceptor or a Connector), it typically does not interact with its initializer.

The following participant is unique to the Acceptor pattern:

- **Acceptor**

- Implements the strategy for passively initializing a Service Handler, which communicates with the peer that actively initiated the connection. The Reactor calls back to the Acceptor’s `accept` method when a connection arrives on the passive-mode `peer_acceptor_endpoint`. The `accept` method uses this passive-mode endpoint to accept connections into the Service Handler’s `peer_stream` and then activate a Service Handler by calling its `open` method.

The following participant is unique to the Connector pattern:

- **Connector**

- Connects and activates a Service Handler. The `connect` method of a Connector implements the strategy for actively initializing a Service Handler, which communicates with the peer that passively accepts the connection. The Connector activates a connected Service Handler by calling its `open` method when initialization is complete. The `complete` method finishes activating Service Handlers whose connections were initiated and completed asynchronously. In this case, the Reactor calls back

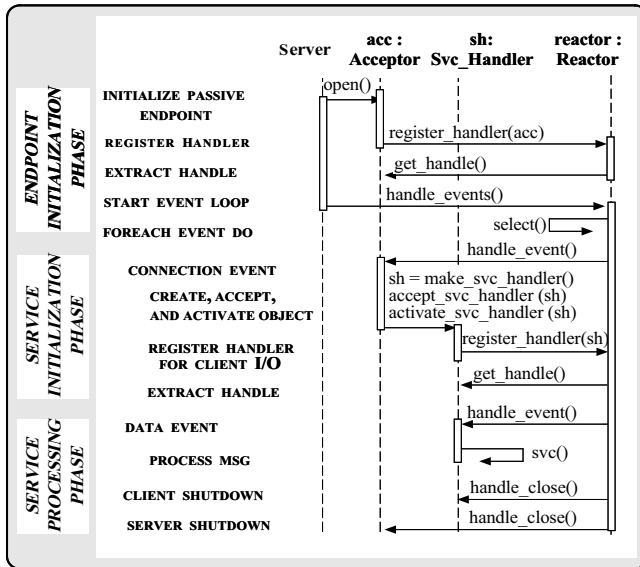


Figure 5: Collaborations Among Participants in the Acceptor Pattern

the `complete` method automatically when an asynchronous connection is established.

4.5 Collaborations

The following section describes the collaborations between participants in the Acceptor and Connector patterns.

4.5.1 Acceptor Collaborations

Figure 5 illustrates the collaboration between participants in the Acceptor pattern. These collaborations are divided into three phases:

- 1. Endpoint initialization phase:** which creates a passive-mode endpoint that is bound to a network address (such as an IP address and port number). The passive-mode endpoint listens for connection requests from peers.
- 2. Service initialization phase:** which activates the Service Handler. When a connection arrives the Reactor calls back to the Acceptor's `accept` method. This method performs the strategy for initializing a Service Handler. This strategy involves assembling the resources necessary to create a new Concrete Service Handler object, accept the connection into this object, and activate the Service Handler by calling its `open` method. The `open` method of the Service Handler then performs service-specific initialization.
- 3. Service processing phase:** Once the connection has been established passively and the service has been initialized, the application enters into a *service processing phase*. This phase performs application-specific tasks that process the data exchanged between the Service Handler and its connected Peer.

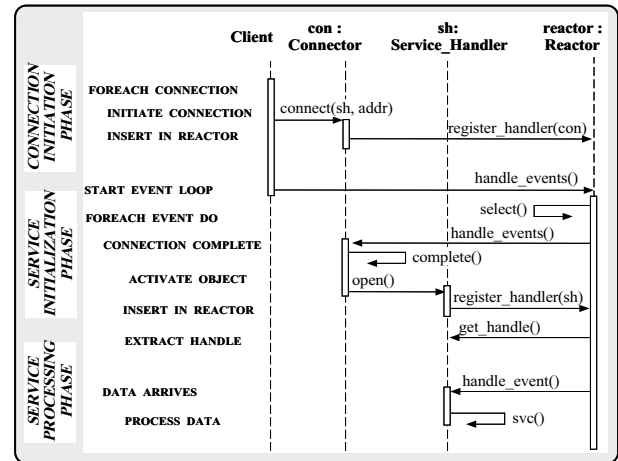


Figure 6: Collaborations Among Connection Pattern Participants for Asynchronous Initialization

4.5.2 Connector Collaborations

The collaborations among participants in the Connector pattern are divided into three phases:

- 1. Connection initiation phase:** which actively connects one or more Service Handlers with their peers. Connections can be initiated synchronously or asynchronously. The Connector's `connect` method implements the strategy for actively establishing connections.
- 2. Service initialization phase:** which activates the Service Handler by calling its `open` method when the connection completes successfully. The `open` method of the Service Handler then performs service-specific initialization.
- 3. Service processing phase:** once the Service Handler is activated, it performs the application-specific service processing using the data exchanged with its connected Peer.

Figure 6 illustrates these three phases of collaboration using *asynchronous* service initialization. Note how the connection initiation phase is temporally separated from the service initialization phase. This enables multiple connection initiations to proceed in parallel within a single thread of control.

The collaboration for *synchronous* service initialization is shown in Figure 7. In this case, the Connector combines the connection initiation and service initialization phases into a single blocking operation.

In general, synchronous connection establishment is useful for the following situations:

- If the latency for establishing a connection is very low (e.g., establishing a connection with a server on the same host via the loopback device);
- If multiple threads of control are available and it is feasible to use a different thread to connect each Service Handler synchronously;

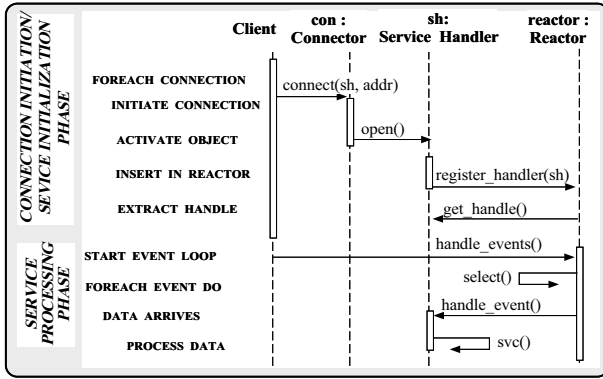


Figure 7: Collaborations Among Connector Pattern Participants for Synchronous Initialization

- If a client application cannot perform useful work until a connection is established.

In contrast, asynchronous connection establishment is useful for the following situations:

- If the connection latency is high and there are many peers to connect with (e.g., establishing a large number of connections over a high-latency WAN);
- If only a single thread of control is available (e.g., if the OS platform does not provide application-level threads);
- If the client application must perform additional work (such as refreshing a GUI) while the connection is in the process of being established.

4.6 Consequences

4.6.1 Benefits

The Acceptor and Connector patterns provide the following benefits:

- **Enhances the reusability, portability, and extensibility of connection-oriented software:** by decoupling mechanisms for passively initializing services from the tasks performed by the services. For instance, the application-independent mechanisms in the Acceptor and Connector are reusable components that know how to establish a connection passively and activate its associated Service Handler. In contrast, the Service Handler knows how to perform the application-specific service processing.

This separation of concerns is achieved by decoupling initialization from service handling, thereby allowing each part to evolve independently. The strategy for active initialization can be written once, placed into a class library or framework, and reused via inheritance, object composition, or template instantiation. Thus, the same passive initialization code need not be rewritten for each application. Services, in contrast, may vary according to different application requirements. By parameterizing the Acceptor and Connector with a

Service Handler, the impact of this variation is localized to a single point in the software.

- **Improves application robustness:** Application robustness is improved by strongly decoupling the Service Handler from the Acceptor. This decoupling ensures that the passive-mode `peer_acceptor_` cannot accidentally be used to read or write data. This eliminates a common class of errors that can arise when programming with weakly typed network programming interfaces such as sockets or TLI [12].

- **Efficiently utilize the inherent parallelism in the network and hosts:** By using the asynchronous mechanisms shown in Figure 6, the Connector pattern can actively establish connections with a large number of peers efficiently over long-latency WANs. This is an important property since a large distributed system may have several hundred Peers connected to a single Gateway. One way to connect all these Peers to the Gateway is to use the synchronous mechanisms shown in Figure 7. However, the round trip delay for a 3-way TCP connection handshake over a long-latency WAN (such as a geosynchronous satellite or trans-atlantic fiber cable) may take several seconds per handshake. In this case, synchronous connection mechanisms cause unnecessary delays since the inherent parallelism of the network and computers is underutilized.

4.6.2 Drawbacks

The Acceptor and Connector patterns have the following drawbacks:

- **Additional indirection:** Both the Acceptor and Connector patterns may require additional indirection compared with using the underlying network programming interfaces directly. However, languages that support parameterized types (such as C++, Ada, or Eiffel), can often implement these patterns with no significant overhead since compilers can inline method calls used to implement the patterns.

- **Additional complexity:** This pattern may add unnecessary complexity for simple client applications that connect with a single server and perform a single service using a single network programming interface.

4.7 Implementation

This section describes how to implement the Acceptor and Connector patterns in C++. The implementation described below is based on reusable components provided in the ACE OO network programming toolkit [4].

Figure 8 divides participants in the Acceptor and Connector patterns into the *Reactive*, *Connection*, and *Application* layers.³ The Reactive and Connection layers perform generic, application-independent strategies for handling

³This diagram illustrates additional Booch notation: directed edges indicate inheritance relationships between classes; a dashed directed edge indicates template instantiation; and a solid circle illustrates a composition relationship between two classes.

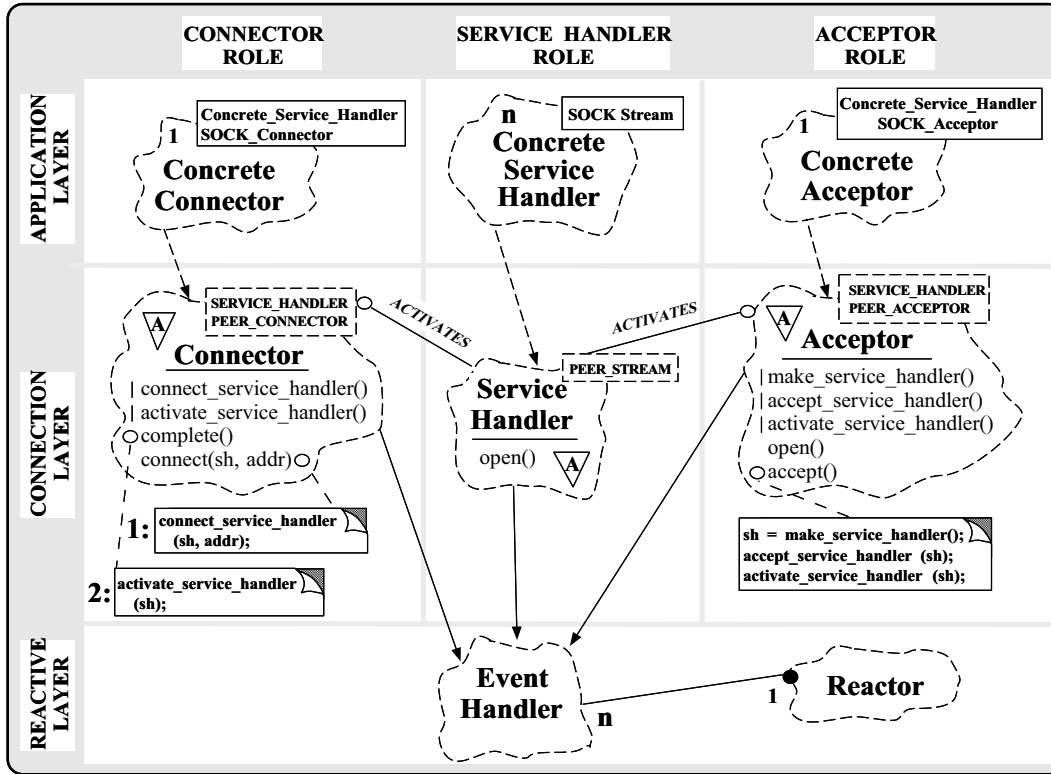


Figure 8: Layering and Partitioning of Participants in the Acceptor and Connector Patterns

events and initializing services, respectively. The Application layer instantiates these generic strategies by providing concrete template classes that establish connections and perform service processing. This separation of concerns increases the reusability, portability, and extensibility in this implementation of the Acceptor and Connector patterns.

The implementations of the Acceptor and Connector patterns are structured very similarly. The Reactive layer is identical in both, and the roles of the Service Handler and Concrete Service Handler are also similar. Moreover, the Acceptor and Concrete Acceptor play roles equivalent to the Connector and Concrete Connector classes. The primary difference between the two patterns is that in the Acceptor pattern these two classes play an *passive* role in establishing a connection. In contrast, in the Connector pattern they play an *active* role.

4.7.1 Reactive Layer

The Reactive layer is responsible for handling events that occur on endpoints of communication represented by I/O handles (also known as “descriptors”). The two participants at this layer, the Reactor and Event Handler, are reused from the Reactor pattern [1]. This pattern encapsulates OS event demultiplexing system calls (such as `select`, `poll` [9], and `WaitForMultipleObjects` [14]) with an extensible and portable callback-driven object-oriented interface. The Reactor pattern enables efficient demultiplexing of multiple types of events from multiple sources within a single

thread of control. The implementation of the Reactor pattern is described in [1]. The two main roles in the Reactive layer are summarized below:

- **Reactor:** This class defines an interface for registering, removing, and dispatching Event Handler objects (such as the Acceptor, Connector, and Service Handler). An implementation of the Reactor interface provides a set of application-independent mechanisms that perform event demultiplexing and dispatching of application-specific Event Handlers in response to events.
- **Event Handler:** This class specifies an interface that the Reactor uses to dispatch callback methods defined by objects that are pre-registered to handle events. These events signify conditions such as a new connection request, a completion of a connection request started asynchronously, or the arrival of data from a connected peer.

4.7.2 Connection Layer

The Connection layer is responsible for (1) creating a Service Handler, (2) passively or actively connecting it with a peer, and (3) activating it once it is connected. Since all behavior in this layer is completely generic, these classes delegate to the concrete IPC mechanism and Concrete Service Handler instantiated by the Application layer (described below). Likewise, the Connection layer delegates to the Reactor to handle initialization-related events (such

as establishing connections asynchronously without requiring multi-threading). The three primary roles in the Connection layer are described below.

• **Service Handler:** This abstract class provides a generic interface for processing services. Applications must customize this class to perform a particular type of service. The following is the interface of the Service Handler:

```
// PEER_STREAM is the type of the
// Concrete IPC mechanism.
template <class PEER_STREAM>
class Service_Handler : public Event_Handler
{
public:
    // Pure virtual method (defined by a subclass).
    virtual int open (void) = 0;

    // Conversion operator needed by
    // Acceptor and Connector.
    operator PEER_STREAM &() { return stream_; }

protected:
    // Concrete IPC mechanism instance.
    PEER_STREAM peer_stream_;
};
```

The open method of a Service Handler is called by the Acceptor or Connector once a connection is established. The behavior of this pure virtual method must be defined by a subclass, which typically performs any service-specific initializations. Service Handler subclasses can also define the service's concurrency strategy. For example, a Service Handler may employ the Reactor [1] pattern to process data from peers in a single-thread of control. Conversely, a Service Handler might use the Active Object pattern [2] to process incoming data in a different thread of control than the one the Acceptor object used to connect it. Section 4.8 illustrates how several different concurrency strategies can be configured flexibly without affecting the structure or behavior of the Acceptor or Connector patterns.

• **Connector:** This abstract class implements the generic strategy for actively initializing communication services. The following class interface illustrates the key methods and objects in the Connector factory:

```
// The SERVICE_HANDLER is the type of service.
// The PEER_CONNECTOR is the type of concrete
// IPC active connection mechanism.
template <class SERVICE_HANDLER,
          class PEER_CONNECTOR>
class Connector : public Event_Handler
{
public:
    enum Connect_Mode {
        SYNC, // Initiate connection synchronously.
        ASYNC // Initiate connection asynchronously.
    };

    // Initialization method.
    Connector (void);

    // Actively connecting and activate a service.
    int connect (SERVICE_HANDLER *sh,
                const PEER_CONNECTOR::PEER_ADDR &addr,
                Connect_Mode mode);

    // Defines the active connection strategy.
    virtual int connect_service_handler
```

```
(SERVICE_HANDLER *sh,
 const PEER_CONNECTOR::PEER_ADDR &addr,
 Connect_Mode mode);

    // Register the SERVICE_HANDLER so that it can
    // be activated when the connection completes.
    int register_handler (SERVICE_HANDLER *sh,
                          Connect_Mode mode);

    // Defines the handler's concurrency strategy.
    virtual int activate_service_handler
        (SERVICE_HANDLER *sh);

    // Activate a SERVICE_HANDLER whose
    // non-blocking connection completed.
    virtual int complete (HANDLE handle);

protected:
    // IPC mechanism that establishes
    // connections actively.
    PEER_CONNECTOR connector_;

    // Collection that maps HANDLES
    // to SERVICE_HANDLER *s.
    Map_Manager<HANDLE, SERVICE_HANDLER *>
        handler_map_;

    // Inherited from the Event_Handler -- will be
    // called back by Eactor when events complete
    // asynchronously.
    virtual int handle_event (HANDLE, EVENT_TYPE);
};

// Useful "short-hand" macros used below.
#define SH SERVICE_HANDLER
#define PC PEER_CONNECTION
```

The Connector is parameterized by a particular type of PEER_CONNECTOR and SERVICE_HANDLER. The PEER_CONNECTOR provides the transport mechanism used by the Connector to actively establish the connection synchronously or asynchronously. The SERVICE_HANDLER provides the service that processes data exchanged with its connected peer. Parameterized types are used to decouple the connection establishment strategy from the type of service handler, network programming interface, and transport layer connection acceptance protocol.

The use of parameterized types helps improve portability by allowing the wholesale replacement of the mechanisms used by the Connector. This makes the connection establishment code portable across platforms that contain different network programming interfaces (such as sockets but not TLI, or vice versa). For example, the PEER_CONNECTOR template argument can be instantiated with either a SOCK Connector or a TLI Connector, depending on whether the platform supports sockets or TLI.

An even more dynamic type of decoupling could be achieved via inheritance and polymorphism by using the Factory Method and Strategy patterns described in [11]. Parameterized types improve run-time efficiency at the expense of additional space and time overhead during program compiling and linking.

The implementation of the Connector's methods is presented below.⁴ The main entry point for a Connector is the connect method:

⁴To save space, most of the error handling in this paper has been omitted.

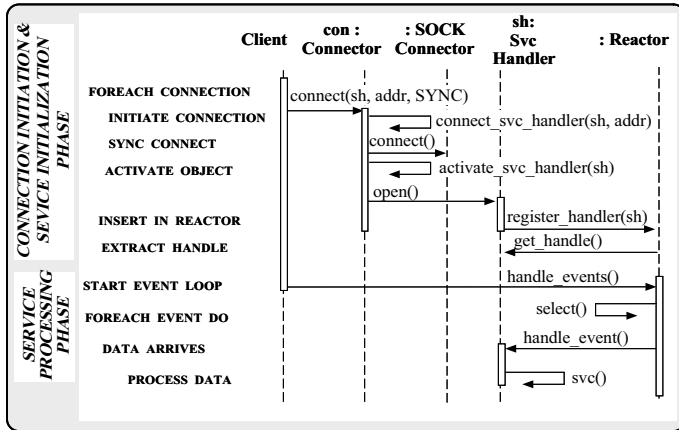


Figure 9: Collaborations Among Participants for Synchronous Connections

```
template <class SH, class PC> int
Connector<SH, PC>::connect
(SERVICE_HANDLER *service_handler,
 const PEER_CONNECTOR::PEER_ADDR &addr,
 Connect_Mode mode)
{
    connect_service_handler (service_handler, addr, mode);
}
```

This method provides the external entry point into the Connector factory. It uses the Bridge pattern⁵ to delegate to the Connector's connection strategy, `connect_service_handler`, which initiates a connection:

```
template <class SH, class PC> int
Connector<SH, PC>::connect_service_handler
(SERVICE_HANDLER *service_handler,
 const PEER_CONNECTOR::PEER_ADDR &remote_addr,
 Connect_Mode mode)
{
    // Delegate to concrete PEER_CONNECTOR
    // to establish the connection.

    if (connector_.connect (*service_handler,
                           remote_addr,
                           mode) == -1) {
        if (mode == ASYNC && errno == EWOULDBLOCK)
            // If the connection hasn't completed and
            // we are using non-blocking semantics then
            // register ourselves with the Reactor
            // Singleton so that it will callback when
            // the connection is complete.
            Reactor::instance ()->register_handler
                (this, WRITE_MASK);

        // Store the SERVICE_HANDLER in the map of
        // pending connections.
        handler_map_.bind
            (connector_.get_handle (), service_handler);
    }
    else if (mode == SYNC)
        // Activate if we connect synchronously.
        activate_service_handler (service_handler);
}
```

⁵The use of the Bridge pattern allows subclasses of Connector to transparently modify the connection strategy, without changing the interface.

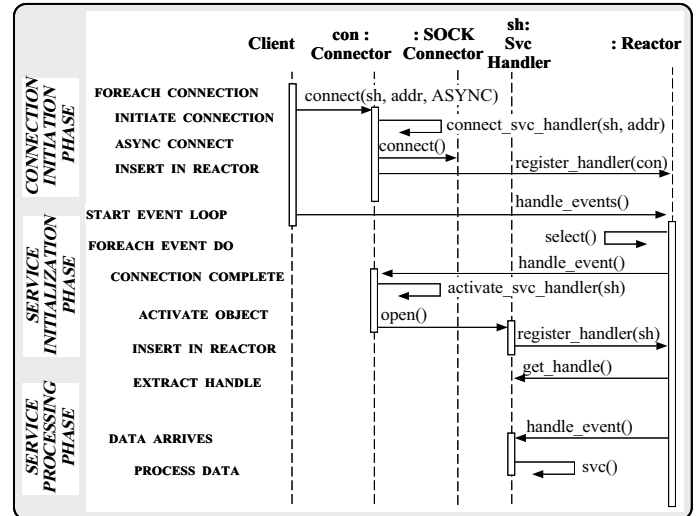


Figure 10: Collaborations Among Participants for Asynchronous Connections

If the value of the `Connect_Mode` parameter is `SYNC` the `SERVICE HANDLER` will be activated after the connection completes synchronously, as illustrated in Figure 9.

To connect with multiple Peers efficiently, however, the Connector must be able to actively establish connections asynchronously, *i.e.*, without blocking the caller. Asynchronous behavior is specified by passing the `ASYNC` connection mode to `Connector::connect`, as illustrated in Figure 10.

When instantiated, the `PEER_CONNECTOR` class provides the concrete IPC mechanism for initiating connections asynchronously. The implementation of the Connector pattern shown here uses asynchronous I/O mechanisms provided by the operating system and communication protocol stack (*e.g.*, by setting sockets into non-blocking mode and using an event demultiplexer like `select` or `WaitForMultipleObjects` to determine when the I/O completes).

The Connector maintains a map of `Service Handlers` whose asynchronous connections are pending completion. Since the Connector inherits from `Event Handler`, the Reactor can automatically call back to the Connector's `handle_event` method when a connection completes. This method is an Adapter that transforms the Reactor's event handling interface to a call to the Connector pattern's `complete` method, as follows:

```
template <class SH, class PC> int
Connector<SH, PC>::handle_event (HANDLE handle,
                                EVENT_TYPE)
{
    // Adapt the Reactor's event handling API to
    // the Connector's API.
    complete (handle);
}
```

The `complete` method then activates a `SERVICE HANDLER` whose non-blocking connection just completed

successfully:

```
template <class SH, class PC> int
Connector<SH, PC>::complete (HANDLE handle)
{
    SERVICE_HANDLER *service_handler = 0;

    // Locate the SERVICE_HANDLER corresponding
    // to the HANDLE.
    handler_map_.find (handle, service_handler);

    // Transfer I/O handle to SERVICE_HANDLER *.
    service_handler->set_handle (handle);

    // Remove handle from Reactor.
    Reactor::instance ()->remove_handler
        (handle, WRITE_MASK);

    // Remove handle from the map.
    handler_map_.unbind (handle);

    // Connection is complete, so activate handler.
    activate_service_handler (service_handler);
}
```

The complete method finds and removes the connected SERVICE_HANDLER from its internal map, transfers the I/O HANDLE to the SERVICE_HANDLER, and initializes the service by calling activate_service_handler. This method delegates to the concurrency strategy designated by the SERVICE_HANDLER::open method, as follows:

```
template <class SH, class PC> int
Connector<SH, PC>::activate_service_handler
    (SERVICE_HANDLER *service_handler)
{
    service_handler->open ();
}
```

Note that activate_service_handler is called when a connection is established successfully, regardless of whether connections are established synchronously or asynchronously. This uniformity of behavior makes it possible to write services whose behavior can be decoupled from the manner by which they are connected.

• **Acceptor:** This abstract class implements the generic strategy for passively initializing communication services. The following class interface illustrates the key methods and objects in the Acceptor factory:

```
// The SERVICE_HANDLER is the type of service.
// The PEER_ACCEPTOR is the type of concrete
// IPC passive connection mechanism.
template <class SERVICE_HANDLER,
         class PEER_ACCEPTOR>
class Acceptor : public Event_Handler {
public:
    // Initialize local_addr listener endpoint
    // and register with Reactor Singleton.
    virtual int open
        (const PEER_ACCEPTOR::PEER_ADDR &local_addr);

    // Factory that creates, connects, and
    // activates SERVICE_HANDLER's.
    virtual int accept (void);

protected:
    // Defines the handler's creation strategy.
    virtual SERVICE_HANDLER *
        make_service_handler (void);
}
```

```
// Defines the handler's connection strategy.
virtual int accept_service_handler
    (SERVICE_HANDLER *);

// Defines the handler's concurrency strategy.
virtual int activate_service_handler
    (SERVICE_HANDLER *);

// Demultiplexing hooks inherited from
// Event_Handler -- used by Reactor for callbacks.
virtual HANDLE get_handle (void) const;
virtual int handle_close (void);

// Invoked when connection requests arrive.
virtual int handle_event (HANDLE, EVENT_TYPE);

private:
    // IPC mechanism that establishes
    // connections passively.
    PEER_ACCEPTOR peer_acceptor_;
};

// Useful "short-hand" macros used below.
#define SH SERVICE_HANDLER
#define PA PEER_ACCEPTOR
```

The Acceptor is parameterized by a particular type of PEER_ACCEPTOR and SERVICE_HANDLER. The PEER_ACCEPTOR provides the transport mechanism used by the Acceptor to passively establish the connection. The SERVICE_HANDLER provides the service that processes data exchanged with its connected peer. Parameterized types are used to decouple the connection establishment strategy from the type of service handler, network programming interface, and transport layer connection initiation protocol.

The use of parameterized types helps improve portability by allowing the wholesale replacement of the mechanisms used by the Acceptor. This makes the connection establishment code portable across platforms that contain different network programming interfaces (such as sockets but not TLI, or vice versa). For example, the PEER_ACCEPTOR template argument can be instantiated with either a SOCK Acceptor or a TLI Acceptor, depending on whether the platform supports sockets or TLI. The implementation of the Acceptor's methods is presented below.

Network applications use the open method to initialize an Acceptor. This method is implemented as follows:

```
template <class SH, class PA> int
Acceptor<SH, PA>::open
    (const PEER_ACCEPTOR::PEER_ADDR &local_addr)
{
    // Forward initialization to the PEER_ACCEPTOR.
    peer_acceptor_.open (local_addr);

    // Register with Reactor.
    Reactor::instance ()->register_handler
        (this, READ_MASK);
}
```

The open method is passed the local_addr network address used to listen for connections. It forwards this address to the passive connection acceptance mechanism defined by the PEER_ACCEPTOR. This mechanism initializes the listener endpoint, which advertises its "service access point" (e.g., IP address and port number) to clients interested in connecting with the Acceptor. The behavior of the listener endpoint

is determined by the type of `PEER ACCEPTOR` instantiated by a user. For instance, it can be a C++ wrapper for sockets [9], TLI [10], STREAM pipes [15], etc.

After the listener endpoint has been initialized, the `open` method registers itself with the Reactor Singleton. The Reactor performs a “double dispatch” back to the `Acceptor`’s `get_handle` method to obtain the underlying `HANDLE`, as follows:

```
template <class SH, class PA> HANDLE
Acceptor<SH, PA>::get_handle
{
    return peer_acceptor_.get_handle ();
}
```

The Reactor stores this `HANDLE` internally and uses it to detect and demultiplex incoming connection from clients. Since the `Acceptor` class inherits from `Event Handler`, the Reactor can automatically call back to the `Acceptor`’s `handle_event` method when a connection arrives from a peer. This method is an Adapter that transforms the Reactor’s event handling interface to a call to the `Acceptor`’s `accept` method, as follows:

```
template <class SH, class PA> int
Acceptor<SH, PA>::handle_event (HANDLE,
                                EVENT_TYPE)
{
    // Adapt the Reactor’s event handling API to
    // the Acceptor’s API.
    accept ();
}
```

As shown below, `accept` is a Template Method that implements the `Acceptor` pattern’s passive initialization strategy for creating a new `SERVICE HANDLER`, accepting a connection into it, and activating the service:

```
template <class SH, class PA> int
Acceptor<SH, PA>::accept (void)
{
    // Create a new SERVICE_HANDLER.
    SH *service_handler = make_service_handler ();

    // Accept connection from client.
    accept_service_handler (service_handler);

    // Activate SERVICE_HANDLER.
    activate_service_handler (service_handler);
}
```

This method is very concise since it factors all low-level details into the concrete `SERVICE HANDLER` and `PEER ACCEPTOR` instantiated via parameterized types. Moreover, all of its behavior is performed by virtual functions, which allow subclasses to extend any or all of the `Acceptor`’s strategies.

The `Acceptor`’s default strategy for creating `SERVICE HANDLERS` is defined by the `make_service_handler` method:

```
template <class SH, class PA> SH *
Acceptor<SH, PA>::make_service_handler (void)
{
    return new SH;
}
```

The default behavior uses a “demand strategy,” which creates a new `SERVICE HANDLER` for every new connection. However, subclasses of `Acceptor` can override this strategy to create `SERVICE HANDLERS` using other strategies (such as creating an individual Singleton [11] or dynamically linking the `SERVICE HANDLER` from a shared library).

The `SERVICE HANDLER` connection acceptance strategy used by the `Acceptor` is defined below by the `accept_service_handler` method:

```
template <class SH, class PA> int
Acceptor<SH, PA>::accept_service_handler
    (SH *handler)
{
    peer_acceptor_->accept_ (*handler);
}
```

The default behavior delegates to the `accept` method provided by the `PEER ACCEPTOR`. Subclasses can override the `accept_service_handler` method to perform more sophisticated behavior (such as authenticating the identity of the client to determine whether to accept or reject the connection).

The `Acceptor`’s `SERVICE HANDLER` concurrency strategy is defined by the `activate_service_handler` method:

```
template <class SH, class PA> int
Acceptor<SH, PA>::activate_service_handler
    (SH *handler)
{
    handler->open ();
}
```

The default behavior of this method is to activate the `SERVICE HANDLER` by calling its `open` method. This allows the `SERVICE HANDLER` to select its own concurrency strategy. For instance, if the `SERVICE HANDLER` inherits from `Event Handler` it can register with the Reactor. This allows the Reactor to dispatch the `SERVICE HANDLER`’s `handle_event` method when events occur on its `PEER STREAM` endpoint of communication. Subclasses can override this strategy to do more sophisticated concurrency activations (such as making the `SERVICE HANDLER` an “active object” [2] that processes data using multi-threading or multi-processing).

When an `Acceptor` terminates, either due to errors or due to the entire application shutting down, the Reactor calls the `Acceptor`’s `handle_close` method, which enables it to release any dynamically acquired resources. In this case, the `handle_close` method simply closes the `PEER ACCEPTOR`’s listener endpoint, as follows:

```
template <class SH, class PA> int
Acceptor<SH, PA>::handle_close (void)
{
    return peer_acceptor_.close ();
}
```

4.7.3 Application Layer

The Application Layer is responsible for supplying a concrete interprocess communication (IPC) mechanism and a concrete

service handler. The IPC mechanisms are encapsulated in C++ classes to simplify programming, enhance reuse, and to enable wholesale replacement of IPC mechanisms. For example, the SOCK Acceptor, SOCK Connector, and SOCK Stream classes used in Section 4.8 are part of the SOCK SAP C++ wrapper library for sockets [16]. Likewise, the corresponding TLI_* classes are part of the TLI SAP C++ wrapper library for the Transport Layer Interface [9]. SOCK SAP and TLI SAP encapsulate the stream-oriented semantics of connection-oriented protocols like TCP and SPX with efficient, portable, and type-safe C++ wrappers.

The three main roles in the Application layer are described below.

- **Concrete Service Handler:** This class implements the concrete application-specific service activated by a Concrete Acceptor or a Concrete Connector. A Concrete Service Handler is instantiated with a specific type of C++ IPC wrapper that exchanges data with its connected peer. The sample code examples in Section 4.8 use a SOCK Stream as the underlying data transport delivery mechanism. It is easy to vary the data transfer mechanism, however, by parameterizing the Concrete Service Handler with a different PEER STREAM (such as an SVR4 TLI Stream or a Win32 Named Pipe Stream).
- **Concrete Connector:** This class instantiates the generic Connector factory with concrete parameterized type arguments for SERVICE_HANDLER and PEER_CONNECTOR.
- **Concrete Acceptor:** This class instantiates the generic Acceptor factory with concrete parameterized type arguments for SERVICE_HANDLER and PEER_ACCEPTOR.

In the sample code in Section 4.8, SOCK Connector and SOCK Acceptor are the underlying transport programming interfaces used to establish connections actively and passively, respectively. However, parameterizing the Connector and Acceptor with different mechanisms (such as a TLI Connector or Named Pipe Acceptor) is straightforward since the IPC mechanisms are encapsulated in C++ wrapper classes.

The following section illustrates sample code that instantiates a Concrete Service Handler, Concrete Connector, and Concrete Acceptor to implement the Peers and Gateway described in Section 3. This particular example of the Application layer customizes the generic initialization strategies provided by the Connector and Acceptor components in the Connection layer. Note how the use of templates and dynamic binding permits specific details (such as the underlying network programming interface or the creation strategy) to change flexibly. For instance, no Connector components must change when the concurrency strategy is modified in Section 4.8.1 and Section 4.8.2.

4.8 Sample Code

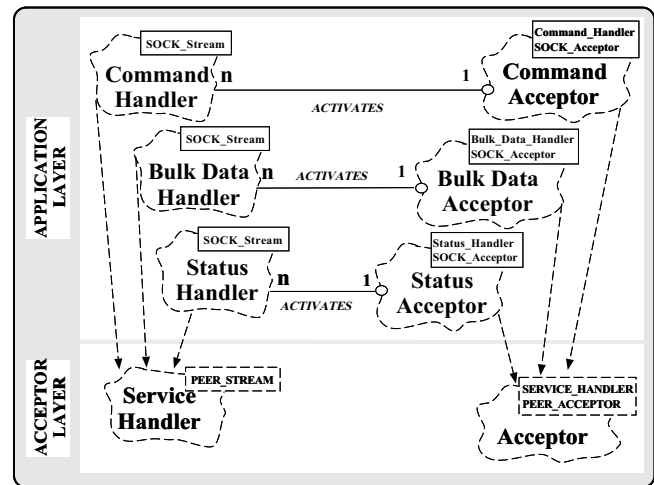


Figure 11: Structure of Acceptor Pattern Participants for Peers

The sample code below illustrates how the Peers and Gateway described in Section 3 use the Acceptor and Connector patterns to simplify the task of passively initializing services. The Peers play the passive role in establishing connections with the Gateway, whose connections are initiated actively by using the Connector pattern. Figure 11 illustrates how participants in the Acceptor pattern are structured in a Peer and Figure 12 illustrates how participants in the Connector pattern are structured in the Gateway.

4.8.1 Peer Components

• **Service Handlers for sending and receiving routing messages in a Peer:** The classes shown below, Status Handler, Bulk Data Handler, and Command Handler, process routing messages sent and received from a Gateway. Since these Concrete Service Handler classes inherit from Service Handler they are capable of being passively initialized by an Acceptor.

To illustrate the flexibility of the Acceptor pattern, each open routine in the Service Handlers implements a different concurrency strategy. In particular, when the Status Handler is activated it runs in a separate thread; the Bulk Data Handler runs as a separate process; and the Command Handler runs in the same thread as the Reactor that demultiplexes connection requests for the Acceptor factories. Note how changes to these concurrency strategies do not affect the implementation of the Acceptor, which is generic and thus highly flexible and reusable.

We start by defining a Service Handler that is specialized for socket-based data transfer:

```
typedef Service_Handler <SOCK_Stream> PEER_HANDLER;
```

The PEER_HANDLER typedef forms the basis for all the subsequent service handlers. For instance, the Status

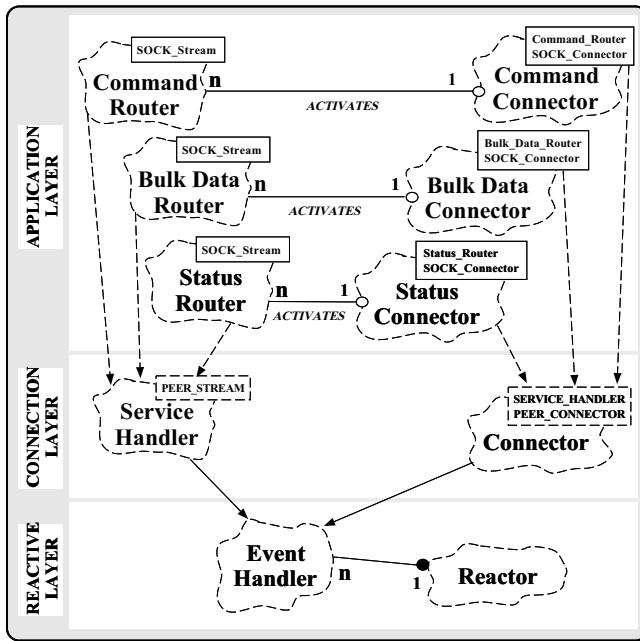


Figure 12: Structure of Connector Pattern Participants for the Gateway

Handler class processes status data sent to and received from a Gateway:

```
class Status_Handler : public PEER_HANDLER
{
public:
    // Performs handler activation.
    virtual int open (void) {
        // Make handler run in separate thread (note
        // that Thread::spawn requires a pointer to
        // a static method as the thread entry point).

        Thread::spawn (&Status_Handler::service_run,
            this);
    }

    // Static entry point into thread, which blocks
    // on the handle_event () call in its own thread.
    static void *service_run (Status_Handler *this_) {
        // This method can block since it
        // runs in its own thread.
        while (this->handle_event () != -1)
            continue;
    }

    // Receive and process status data from Gateway.
    virtual int handle_event (void) {
        char buf[MAX_STATUS_DATA];
        stream_.recv (buf, sizeof buf);
        // ...
    }

    // ...
};
```

The following class processes bulk data sent to and received from the Gateway.

```
class Bulk_Data_Handler : public PEER_HANDLER
{
public:
    // Performs handler activation.
```

```
virtual int open (void) {
    // Handler runs in separate process.
    if (fork () > 0) // In parent process.
        return 0;
    else // In child process.

        // This method can block since it
        // runs in its own process.
        while (handle_event () != -1)
            continue;
}

// Receive and process bulk data from Gateway.
virtual int handle_event (void) {
    char buf[MAX_BULK_DATA];
    stream_.recv (buf, sizeof buf);
    // ...
}

// ...
};
```

The following class processes bulk data sent to and received from a Gateway:

```
class Command_Handler : public PEER_HANDLER
{
public:
    // Performs handler activation.
    virtual int open (void) {
        // Handler runs in same thread as main
        // Reactor singleton.
        Reactor::instance ()->register_handler
            (this, READ_MASK);
    }

    // Receive and process command data from Gateway.
    virtual int handle_event (void) {
        char buf[MAX_COMMAND_DATA];
        // This method cannot block since it borrows
        // the thread of control from the Reactor.
        stream_.recv (buf, sizeof buf);
        // ...
    }

    //...
};
```

• **Acceptors for creating Peer Service Handlers:** The `s_acceptor`, `bd_acceptor`, and `c_acceptor` objects shown below are Concrete Acceptor factories that create and activate Status Handlers, Bulk Data Handlers, and Command Handlers, respectively.

```
// Accept connection requests from Gateway and
// activate Status_Handler.
Acceptor<Status_Handler, SOCK_Acceptor> s_acc;

// Accept connection requests from Gateway and
// activate Bulk_Data_Handler.
Acceptor<Bulk_Data_Handler, SOCK_Acceptor> bd_acc;

// Accept connection requests from Gateway and
// activate Command_Handler.
Acceptor<Command_Handler, SOCK_Acceptor> c_acc;
```

• **The Peer Main function:** The main program initializes the concrete Acceptor factories by calling their open methods with the well-known ports for each service. As shown in Section 4.7.2, the `Acceptor::open` method registers itself with an instance of the Reactor. The program then enters an event loop that uses the Reactor to

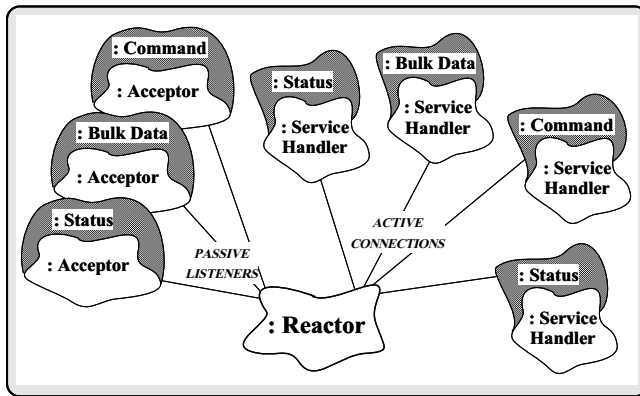


Figure 13: Object Diagram for the Peer Acceptor Pattern

detect connection requests from the Gateway. When connections arrive, the Reactor calls back to the appropriate Acceptor, which creates a new `PEER_HANDLER` to perform the service, accepts the connection into the handler, and activates the handler.

```
// Main program for the Peer.

int main (void)
{
    // Initialize acceptors with their well-known ports.
    s_acc.open (INET_Addr (STATUS_PORT));
    bd_acc.open (INET_Addr (BULK_DATA_PORT));
    c_acc.open (INET_Addr (COMMAND_PORT));

    // Loop forever handling connection request
    // events and processing data from the Gateway.

    for (;;)
        Reactor::instance ()->handle_events ();
}
```

Figure 13 illustrates the relationship between Acceptor pattern objects in the Peer after four connections have been established. While the various Handlers exchange data with the Gateway, the Acceptors continue to listen for new connections.

4.8.2 Gateway Components

• **Service Handlers for Gateway routing:** The classes shown below, Status Router, Bulk Data Router, and Command Router, route data they receive from a source Peer to one or more destination Peers. Since these Concrete Service Handler classes inherit from Service Handler they can be actively connected and initialized by a Connector.

To illustrate the flexibility of the Connector pattern, each open routine in a Service Handler implements a different concurrency strategy. In particular, when the Status Router is activated it runs in a separate thread; the Bulk Data Router runs as a separate process; and the Command Router runs in the same thread as the Reactor that demultiplexes connection completion events for the Connector factory. As with the Acceptor, note

how changes to these concurrency strategies do not affect the implementation of the Connector, which is generic and thus highly flexible and reusable.

We'll start by defining a Service Handler that is specialized for socket-based data transfer:

```
typedef Service_Handler <SOCK_Stream> PEER_ROUTER;
```

This class forms the basis for all the subsequent routing services. For instance, the Status Router class routes status data from/to Peers:

```
class Status_Router : public PEER_ROUTER
{
public:
    // Activate router in separate thread.
    virtual int open (void) {
        // Thread::spawn requires a pointer to a
        // static method as the thread entry point).
        Thread::spawn (&Status_Router::service_run,
            this);
    }

    // Static entry point into thread, which blocks
    // on the handle_event() call in its own thread.
    static void *service_run (Status_Router *this_) {
        // This method can block since it
        // runs in its own thread.
        while (this_->handle_event () != -1)
            continue;
    }

    // Receive and route status data from/to Peers.
    virtual int handle_event (void) {
        char buf[MAX_STATUS_DATA];
        peer_stream_.recv (buf, sizeof buf);
        // Routing takes place here...
    }

    // ...
};
```

The Bulk Data Router routes bulk data from/to Peers:

```
class Bulk_Data_Router : public PEER_ROUTER
{
public:
    // Activates router in separate process.
    virtual int open (void) {
        if (fork () > 0) // In parent process.
            return 0;
        else // In child process.

            // This method can block since it
            // runs in its own process.
            while (handle_event () != -1)
                continue;
    }

    // Receive and route bulk data from/to Peers.
    virtual int handle_event (void) {
        char buf[MAX_BULK_DATA];
        peer_stream_.recv (buf, sizeof buf);
        // Routing takes place here...
    }
};
```

The Command Router class routes Command data from/to Peers:

```
class Command_Router : public PEER_ROUTER
{
```

```

public:
    // Activates router in same thread as Connector.
    virtual int open (void) {
        Reactor::instance ()->register_handler
            (this, READ_MASK);
    }

    // Receive and route command data from/to Peers.
    virtual int handle_event (void) {
        char buf[MAX_COMMAND_DATA];
        // This method cannot block since it borrows
        // the thread of control from the Reactor.
        peer_stream_.recv (buf, sizeof buf);
        // Routing takes place here...
    }
};

```

• **The Gateway Main function:** The main program for the Gateway is shown below. The `get_peer_addrs` function creates the Status, Bulk Data, and Command Routers that route messages through the Gateway. This function (whose implementation is not shown) reads a list of Peer addresses from a configuration file. Each Peer address consists of an IP address and a port number. Once the Routers are initialized, the Connector factories defined above initiate all the connections asynchronously (indicated by passing the ASYNC flag to the connect method).

```

// Main program for the Gateway.

// Define a Connector factory specialized for
// PEER_ROUTERS.

typedef Connector<PEER_ROUTERS, SOCK_Connector>
    PEER_CONNECTOR;

// Obtain lists of Status_Routers,
// Bulk_Data_Routers, and Command_Routers
// from a config file.

void get_peer_addrs (Set<PEER_ROUTERS> &peers);

int main (void)
{
    // Connection factory for PEER_ROUTERS.
    PEER_CONNECTOR peer_connector;

    // A set of PEER_ROUTERS that perform
    // the Gateway's routing services.
    Set<PEER_ROUTER> peers;
    PEER_ROUTER *peer;

    // Get set of Peers to connect with.
    get_peer_addrs (peers);

    // Iterate through all the Routers and
    // initiate connections asynchronously.

    for (Set_Iter<PEER_ROUTER> set_iter (peers);
         set_iter.next (peer) != 0;
         set_iter++)
        peer_connector.connect (peer,
                                peer->address (),
                                PEER_CONNECTOR::ASYNC);

    // Loop forever handling connection completion
    // events and routing data from Peers.

    for (;;)
        Reactor::instance ()->handle_events ();

    /* NOTREACHED */
    return 0;
}

```

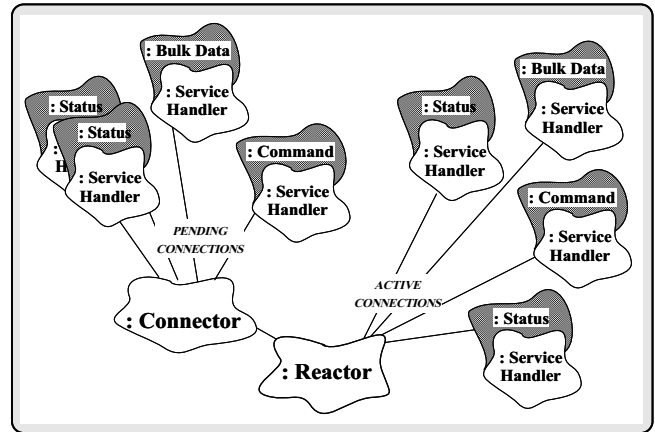


Figure 14: Object Diagram for the Peer Connector Pattern

All connections are invoked asynchronously. They complete concurrently via `Connector::complete` methods that are called back within the Reactor's event loop. The Reactor also demultiplexes and dispatches routing events for Command Router objects, which run in the Reactor's thread of control. The Status Routers and Bulk Data Routers execute in separate threads and processes, respectively.

Figure 14 illustrates the relationship between objects in the Gateway after four connections have been established. Four other connections that have not yet completed are owned by the Connector. When all Peer connections are completely established, the Gateway can route and forward messages sent to it by Peers.

4.9 Known Uses

The Acceptor and Connector patterns have been used in the following frameworks, toolkits, and systems:

- **UNIX network superservers:** such as `inetd` [9], `listen` [10], and the Service Configurator daemon from the ASX framework [4]. These superservers utilize a master Acceptor process that listens for connections on a set of communication ports. Each port is associated with a communication-related service (such as the standard Internet services `ftp`, `telnet`, `daytime`, and `echo`). The Acceptor pattern decouples the functionality in the `inetd` superserver into two separate parts: one for establishing connections and another for receiving and processing requests from peers. When a service request arrives on a monitored port, the Acceptor process accepts the request and dispatches an appropriate pre-registered handler to perform the service.

- **Ericsson EOS Call Center Management System:** this system uses the Acceptor and Connector patterns to allow application-level Call Center Manager Event Servers [17] to actively establish connections with passive Supervisor in a distributed system.

• **Project Spectrum:** The high-speed medical image transfer subsystem of project Spectrum [18] uses the Acceptor and Connector patterns to passively establish connections and initialize application services for storing large medical images. Once connections are established, applications then send and receive multi-megabyte medical images to and from these image stores.

• **ACE Framework:** Implementations of the Reactor, Service Handler, Connector, and Acceptor classes described in this paper are all provided as reusable components in the ACE object-oriented network programming framework [4].

4.10 Related Patterns

The Acceptor and Connector patterns use the Template Method and Factory Method patterns [11]. The Acceptor's accept and the Connector's connect and complete functions are Template Methods that implements a generic service initialization Strategy for connecting with peers and activating a Service Handler when the connections is established. The use of the Template Method pattern allows subclasses to modify the specific details of creating, connecting, and activating Service Handlers. The Factory Method pattern is used to decouple the creation of a Service Handler from its subsequent use.

The Connector pattern has an intent similar to the Client-Dispatcher-Server pattern described in [19]. They both are concerned with separating active connection establishment from the subsequent service. The primary difference is that the Connector pattern addresses both synchronous and asynchronous service initialization, whereas the Client-Dispatcher-Server pattern focuses on synchronous connection establishment.

5 Concluding Remarks

This paper describes the Acceptor and Connector patterns and gives a detailed example illustrating how to use them. Implementations of the Acceptor, Connector, and Reactor patterns described in this paper are freely available via the World Wide Web at URL <http://www.cs.wustl.edu/~schmidt/ACE.html>. This distribution contains complete source code, documentation, and example test drivers for the C++ components developed as part of the ACE object-oriented network programming toolkit [4] developed at the University of California, Irvine and Washington University, St. Louis. The ACE toolkit is currently being used on communication software at many companies including Bellcore, Siemens, DEC, Motorola, Ericsson, and Kodak.

References

- [1] D. C. Schmidt, "Reactor: An Object Behavioral Pattern for Concurrent Event Demultiplexing and Event Handler Dispatching," in *Pattern Languages of Program Design* (J. O. Coplien and D. C. Schmidt, eds.), Reading, MA: Addison-Wesley, 1995.
- [2] R. G. Lavender and D. C. Schmidt, "Active Object: an Object Behavioral Pattern for Concurrent Programming," in *Pattern Languages of Program Design* (J. O. Coplien, J. Vlissides, and N. Kerth, eds.), (Reading, MA), Addison-Wesley, 1996.
- [3] P. Jain and D. C. Schmidt, "Service Configurator: A Pattern for Dynamic Configuration and Reconfiguration of Communication Services," in *Submitted to the 3rd Pattern Languages of Programming Conference*, September 1996.
- [4] D. C. Schmidt, "ACE: an Object-Oriented Framework for Developing Distributed Applications," in *Proceedings of the 6th USENIX C++ Technical Conference*, (Cambridge, Massachusetts), USENIX Association, April 1994.
- [5] I. Pyarali, T. H. Harrison, and D. C. Schmidt, "Design and Performance of an Object-Oriented Framework for High-Performance Electronic Medical Imaging," in *Proceedings of the 2nd Conference on Object-Oriented Technologies and Systems*, (Toronto, Canada), USENIX, June 1996.
- [6] J. Postel, "Transmission Control Protocol," *Network Information Center RFC 793*, pp. 1–85, Sept. 1981.
- [7] S. J. Leffler, M. McKusick, M. Karels, and J. Quarterman, *The Design and Implementation of the 4.3BSD UNIX Operating System*. Addison-Wesley, 1989.
- [8] W. R. Stevens, *TCP/IP Illustrated, Volume 1*. Reading, Massachusetts: Addison Wesley, 1993.
- [9] W. R. Stevens, *UNIX Network Programming*. Englewood Cliffs, NJ: Prentice Hall, 1990.
- [10] S. Rago, *UNIX System V Network Programming*. Reading, MA: Addison-Wesley, 1993.
- [11] E. Gamma, R. Helm, R. Johnson, and J. Vlissides, *Design Patterns: Elements of Reusable Object-Oriented Software*. Reading, MA: Addison-Wesley, 1995.
- [12] D. C. Schmidt, T. H. Harrison, and E. Al-Shaer, "Object-Oriented Components for High-speed Network Programming," in *Proceedings of the 1st Conference on Object-Oriented Technologies and Systems*, (Monterey, CA), USENIX, June 1995.
- [13] G. Booch, *Object Oriented Analysis and Design with Applications (2nd Edition)*. Redwood City, California: Benjamin/Cummings, 1993.
- [14] H. Custer, *Inside Windows NT*. Redmond, Washington: Microsoft Press, 1993.
- [15] D. L. Presotto and D. M. Ritchie, "Interprocess Communication in the Ninth Edition UNIX System," *UNIX Research System Papers, Tenth Edition*, vol. 2, no. 8, pp. 523–530, 1990.
- [16] D. C. Schmidt, "IPC_SAP: An Object-Oriented Interface to Interprocess Communication Services," *C++ Report*, vol. 4, November/December 1992.
- [17] D. C. Schmidt and T. Suda, "An Object-Oriented Framework for Dynamically Configuring Extensible Distributed Communication Systems," *IEEE/BCS Distributed Systems Engineering Journal (Special Issue on Configurable Distributed Systems)*, vol. 2, pp. 280–293, December 1994.
- [18] G. Blaine, M. Boyd, and S. Crider, "Project Spectrum: Scalable Bandwidth for the BJC Health System," *HIMSS, Health Care Communications*, pp. 71–81, 1994.
- [19] F. Buschmann, R. Meunier, H. Rohnert, P. Sommerlad, and M. Stal, *Pattern-Oriented Software Architecture - A System of Patterns*. Wiley and Sons, 1996.